

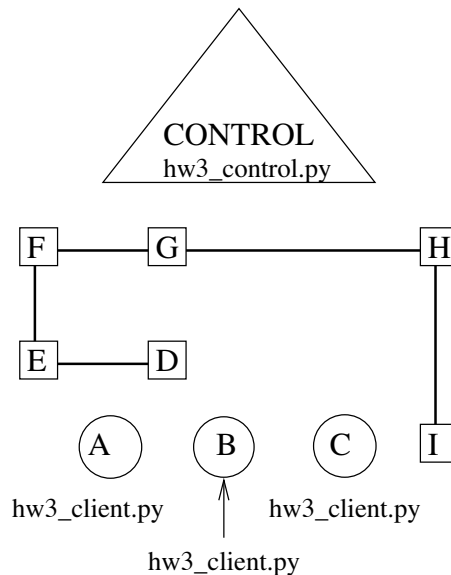
CSCI 4220 Assignment 3

Assignment 3: Mobile Sensor Network Relay

Overview

Consider a mobile sensor application running over TCP with the topology (layout) below. The figure is drawn to scale. Each of the circles (A,B,C) represents a “sensor” which can move around in 2-D space (x and y coordinates) and runs your client application, `hw3_client.py` or `client.out`. The squares represent “base stations” (D,E,F,G,H,I) which do not run any programs and do not move. Finally, the large triangle (CONTROL) is an abstract “control server” (it has no position and can reach all base stations and all clients) that will run your server application `hw3_control.py` or `server.out`. The control server will be used to relay TCP messages between sensors (which never directly communicate) and to manage messages that are moving between base stations (since base stations do not have any code and do not use sockets).

Every base station or sensor has a unique ID consisting of only letters, underscores, and numbers. The control server will always have an ID of CONTROL. You can assume we will not provide invalid IDs or duplicate IDs.



Every sensor has a maximum communication range, and sensors can only send or receive messages if **both** the sender and receiver can reach each other (have communication ranges greater than or equal to the distance between sender and receiver). However, instead of directly sending to or receiving from another sensor/base station, they will always communicate to/from the server. If a base station wants to send a message to a sensor, then the control server will be the sender and the sensor will be the receiver. Base stations can only communicate with other base stations that are connected - this is represented through a “base station file” which is described later in this handout. When communicating with sensors, base stations are assumed to have a range of infinity.

If a base station or sensor has a message it wants to send, it decides where to send it based on a three step decision process. First if the intended recipient (who the message is addressed to) is in range (meaning both

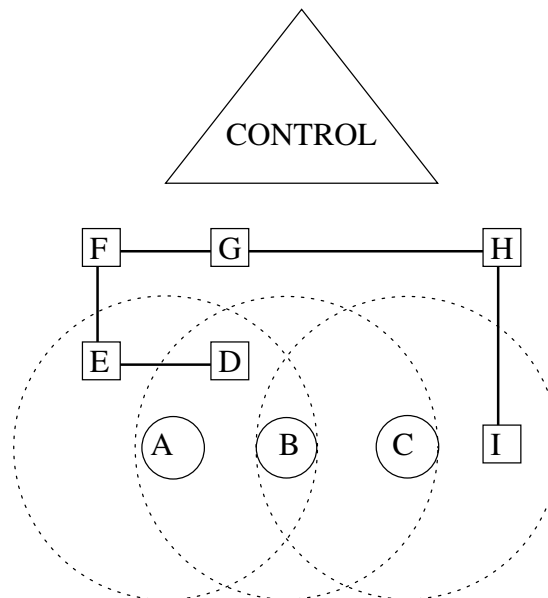
the sender and destination can reach each other directly), then the message is sent to the recipient. Otherwise, the sender sorts all in-range sensors and base stations by distance to the destination (using the “Euclidean distance”). The sender then sends the message to the base station or sensor with the lowest distance to the intended recipient (assuming this would not cause a cycle, see below for an example), resolving ties by choosing the lexicographically smaller ID.

You can assume that we will not give any incorrectly formatted input, and we will not give any IDs that do not exist. You can also assume that clients will not disconnect until the end of the program. You may use any of the techniques we have discussed in this class so far. We will briefly discuss `select()` and socket programming in Python in class, along with an overview of the assignment. Coming from C the concepts are the same, but it’s a lot cleaner. You are not required to use `select()`, however we will not be discussing threading in Python as a class so if you want to take a threaded approach you will have to do your own research.

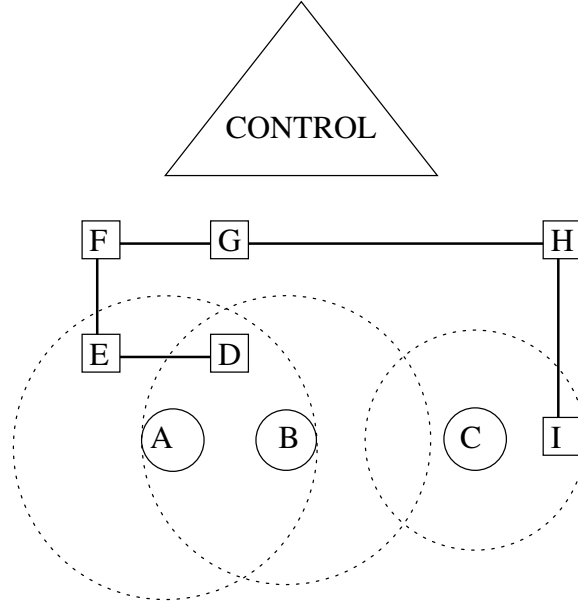
Invocation of the program will use hostnames. The tests will be inputting hostnames instead of IP addresses. To save you a few includes, we have provided `hw3_starter.py` under Course Materials which shows how to get your local hostname and IP address and resolve a remote hostname into an IP address. Basically you will need to make use of `gethostbyname()` to translate those to IP addresses.

In summary, the control server listens for incoming connections, processes commands given via standard input, and reacts to messages from sensors. (Base station to base station communication can be done without any sockets since it’s all in the same `hw3_control` process) The sensors running (`hw3_client.py`) process commands given via standard input, and react to messages received from the control server.

Examples of Communication



In the diagram shown above, a message from B to C can be directly delivered (so it will be sent from B to CONTROL to C). A message from A to C cannot be directly delivered, since that would require A and C being in range of each other. Instead, A will pass the message to B, and B will pass it to C. Recall that sensors do not directly communicate with each other, so the message will actually travel from A to CONTROL to B to CONTROL to C.



The diagram above pictures a slightly different topology in which even **B** cannot directly reach **C**. If **B** wants to send a message to **C**, it will first send to **D** (which is the reachable sensor/base station that is closest to **C**). **D** will send to **A** (sending back to **B** would cause a cycle, and **A** is slightly closer than **E** is to **C**). **A** will have no choice but to send to **E** (anything else would cause a cycle). **E** will send to **F** (also because anything else will cause a cycle), and the remaining steps follow similarly: **F** to **G** to **H** to **I** to **C**.

However, because sensors do not directly communicate, the actual path is:

B to **CONTROL** (handling **D**) to **A** to **CONTROL** (handling **E** to **F** to **G** to **H** to **I**, no socket communication needed!) to **C**.

Base Station File Format

The base station file contains the following information in each row (ending in a newline character `\n`), with each field separated by a space:

- `[BaseID]` `[XPos]` `[YPos]` `[NumLinks]` `[ListOfLinks]`
- `[BaseID]` is the unique string identifier for the base station, and only contains letters, numbers and underscores.
- `[XPos]` is a floating point number representing the x-coordinate of this base station
- `[YPos]` is a floating point number representing the y-coordinate of this base station
- `[NumLinks]` is an integer that is the degree of the base station (i.e. how many other base stations it is connected to)
- `[ListOfLinks]` is a space-separated list of base station IDs that this base station is connected to. All connections are undirected, (i.e. BaseA is connected to BaseB iff BaseB is connected to BaseA).

You can assume that from any base station it is possible to reach all other base stations eventually, though it may require using multiple base stations along the way. In graph theory terms, the graph is all one giant component.

Message Formatting

Messages that are being sent to another base station/sensor will look as follows, with each field separated by a space:

- DATAMESSAGE [OriginID] [NextID] [DestinationID] [HopListLength] [HopList]
- [OriginID] is the ID of the base station/sensor that initially sent the data message.
- [NextID] is the ID that the data message is currently being sent to.
- [DestinationID] is the ID of the final destination that the data message is trying to be sent to.
- [HopListLength] is the length of [HopList]
- [HopList] contains the IDs that have been used so far to deliver the data message

Control Server Arguments, Input and Messages

The server is invoked by running: `python3 -u hw3_control.py [control port] [base station file]`

- [control port] is the port the server should listen on for sensors to connect to.
- [base station file] is the name of the file the server should read to get information about base stations and the base station network.

The following commands can be read from standard input:

- QUIT causes the server to clean up any memory and any sockets that are in use, and then terminate.
- If the server receives a WHERE message from a sensor, it should respond with a message in the following format:
THERE [NodeID] [XPosition] [YPosition] where [XPosition] and [YPosition] are the x and y-coordinates of the requested base station/sensor and [NodeID] is its ID. Remember that we are assuming all IDs provided are valid.
- If the server receives an UPDATEPOSITION message from a sensor, it should respond with a message in the following format: REACHABLE [NumReachable] [ReachableList]
where [ReachableList] contains all sensor and base station IDs and positions that are currently reachable by the sensor, and [NumReachable] is the number of IDs in [ReachableList]. Each entry in ReachableList is actually 3 strings: [ID] [XPosition] [YPosition]. Entries are separated by a space.
- If the server receives a data message with a [NextID] that is a sensor's ID, the server is acting as a sensor-to-sensor relay, and it should deliver the message to the destination.

If a base station receives a data message (either the server receives a message, or we are internally passing a message from one base station to another base station):

- If the [DestinationID] matches its [BaseID] it should print the following to standard output:
[BaseID]: Message from [OriginID] to [DestinationID] successfully received.
- Otherwise, if a base station with ID [BaseID] receives a data message (starts with DATAMESSAGE) with a [DestinationID] that is not its own ID, and all reachable sensors/base stations are already in the [HopList], the data message should be discarded and the following should be printed to standard output: [BaseID]: Message from [OriginID] to [DestinationID] could not be delivered.
- Otherwise, the base station should send a new DATAMESSAGE command to the next sensor/base station based on the decision-making described in the "Overview" section. This DATAMESSAGE should look the same, but the [NextID] should reflect the new next hop, [HopList] should also contain the [BaseID] of the base station and the [HopListLength] should be increased by 1. A message should be printed to standard output:
[BaseID]: Message from [OriginID] to [DestinationID] being forwarded through [BaseID]
- If the [OriginID] is the current base station, and the [NextID] and [DestinationID] match, then the base station should print the following message to standard output:

- [BaseID]: Sent a new message directly to [DestinationID].
- Otherwise if the [OriginID] is the current base station, then the base station should print the following message to standard output:
[BaseID]: Sent a new message bound for [DestinationID].

Sensor (Client) Arguments, Input and Messages

Sensors invoke the client code by running: `python3 -u hw3_client.py [control address] [control port] [SensorID] [SensorRange] [InitialXPosition] [InitialYPosition]`

- [control address] is the address or hostname the control server can be reached at
- [control port] is the port the control server will be listening on for
- [SensorID] is the ID of the sensor and only contains letters and numbers.
- [SensorRange] is an integer value for the maximum communication range of the sensor
- [InitialXPosition] is the initial x-coordinate of the sensor
- [InitialYPosition] is the initial y-coordinate of the sensor

When the sensor starts up, before processing any standard input it should first connect to the control server and send an UPDATEPOSITION message (see below). The following commands can be read from standard input:

- MOVE [NewXPosition] [NewYPosition] causes the sensor to update its location to the x-coordinate specified by [NewXPosition] and the y-coordinate specified by [NewYPosition]. The sensor should also send an UPDATEPOSITION command to the server.
- SENDDATA [DestinationID] causes the sensor to generate a new DATAMESSAGE with a destination of [DestinationID]. The sensor should first send an UPDATEPOSITION message to the server to get an up-to-date list of reachable sensors and base stations.
- QUIT causes the client program to clean up any memory and any sockets that are in use, and then terminate.
- Sensors can send a WHERE message to the control server:
WHERE [SensorID/BaseID] to get the location of a particular base station or sensor ID from the control server. It should not take any other actions until it gets a THERE message back from the server.

Sensors can send an UPDATEPOSITION message to the control server:

UPDATEPOSITION [SensorID] [SensorRange] [CurrentXPosition] [CurrentYPosition] where the positions are the current x and y coordinates of the sensor. After sending an UPDATEPOSITION message, the client should not do anything else until it receives a REACHABLE message from the server. When a REACHABLE message is received, the sensor should print the following message to standard output:

[SensorID]: After reading REACHABLE message, I can see: [SortedReachableList]

where [SortedReachableList] is a list of all the sensors/base stations that can be reached, sorted by their [SensorID]/[BaseID]

If a sensor receives a DATAMESSAGE from the control server:

- If the [DestinationID] matches its [SensorID] it should print the following to standard output:
[SensorID]: Message from [OriginID] to [DestinationID] successfully received.
- Otherwise, the sensor should first send an UPDATEPOSITION message to CONTROL (even it hasn't moved) so that it gets an updated list of reachable sensors and base stations. The sensor should next send a WHERE message to the server requesting the current position of the [DestinationID]. The sensor should then send a new DATAMESSAGE command to the next sensor/base station based on the decision-making described in the "Overview" section. This DATAMESSAGE should look the same, but the [NextID] should reflect the new next hop, [HopList] should also contain the [SensorID] of the sensor and the [HopListLength] should be increased by 1. A message should be printed to standard output:
[SensorID]: Message from [OriginID] to [DestinationID] being forwarded through [SensorID]
- But if a sensor with [SensorID] receives a data message with a [DestinationID] that is not its own ID, and all reachable sensors/base stations are already in the [HopList] after UPDATEPOSITION, the message should be discarded and the following message should be printed to standard output:
[SensorID]: Message from [OriginID] to [DestinationID] could not be delivered.
- **IMPORTANT:** Sensors will never directly talk to each other. If a sensor wants to pass a message to another sensor, it will update the [NextID] field in the DATAMESSAGE to be the sensor it wants to pass the message to, and send the message to CONTROL.
- If the [OriginID] is the current sensor, and the [NextID] and [DestinationID] match, then the sensor should print the following message to standard output:
[SensorID]: Sent a new message directly to [DestinationID].
- Otherwise if the [OriginID] is the current sensor, then the sensor should print the following message to standard output:
[SensorID]: Sent a new message bound for [DestinationID].

Submission

You can provide any decisions or other comments for the graders to read in a **README.txt** file. You can do this assignment in Python 3 or in C/C++ (book library provided). The arguments are the same in both cases. You should submit two Python files, **hw3_client.py** and **hw3_control.py** which will be invoked for sensors and the control server respectively, or a **Makefile** (case sensitive) that expects your source files to be in a directory called **part1** – see the provided Makefile.